



UNIVERSITÄT
ZU KÖLN



Computer Vision

Prof. Dr. rer. nat. Katarzyna Bozek



UNIVERSITÄT
ZU KÖLN



Journal Club

Journal Club

- Presentation of a paper in groups of 1-2 people
 - ~30 min presentation, followed by a discussion
 - 1-2 paper per slot
- Everybody should have skimmed over the paper beforehand
- Each student should have presented at least once
- Presentations will be taken into account for the final grade

Journal Club

- 1. Paper: "[The Watershed Transformation Applied to Image Segmentation](#)"
 - Technique used in image processing for segmenting objects
 - Mimics the process of flooding a landscape to identify catchments basins
- 2. Paper: "[Mean shift: a robust approach toward feature space analysis](#)"
 - Clustering algorithm
 - Identifies clusters in data without requiring prior assumptions about their shapes or sizes
 - Used for segmentation
- Presentations will take place on 02.05.2024



UNIVERSITÄT
ZU KÖLN



How to read a paper

How to Read a Paper – S. Keshav, 2007

How to read a paper

- Researchers read a lot of paper
 - Reviews
 - Catching up to state-of-the-art methods
 - Literature survey
- Can be very time consuming if not done in a systematic way
- Three-pass approach
 - Read the paper in (at max.) three passes
 - Number of passes depend the relevance of the pass before

First Pass

- Getting brief overview of the paper (~ 10 min.)
 - Title, abstract, introduction
 - Read the headings
 - Skim over the references
 - Conclusion
- Be able to tell the **Category** and the **Context** of the paper
- Do the assumption appear to be valid for you (**Correctness**)?
- Is it well written (**Clarity**)?
- What are the main **Contributions**?

Continue reading?

- Decide if the paper is still relevant for you
- Are you able to understand it?
- Maybe you decide the assumption of the author is invalid...
- The paper could also be out of your research, but just for this moment

If you decide the paper is relevant, valid and of interest => Second pass

2nd pass

- Second Pass (~ 1 hr.)
 - Reading the paper with greater care
 - But ignore details like proofs
 - Look at the illustrations in the paper (figures, diagrams, graphs)
 - Pay extra attentions to graphs
 - Are the axes labelled correctly?
 - Do the results include error bars?
 - This can be a hint on the quality of the paper and the effort the author put into it
 - Mark relevant unread references for future reading
 - Write down comments on the go

2nd pass

3.1 Encoder and Decoder Stacks

Encoder: The encoder is composed of a stack of $N = 6$ identical layers. Each layer has two sub-layers. The first is a multi-head self-attention mechanism, and the second is a simple, position-wise fully connected feed-forward network. We employ a residual connection [11] around each of the two sub-layers, followed by layer normalization [1]. That is, the output of each sub-layer is $\text{LayerNorm}(x + \text{Sublayer}(x))$, where $\text{Sublayer}(x)$ is the function implemented by the sub-layer itself. To facilitate these residual connections, all sub-layers in the model, as well as the embedding layers, produce outputs of dimension $d_{\text{model}} = 512$.

After the 2nd pass

- You should comprehend the content of the paper
- Be able to summarize the main trust to a colleague (with supporting evidence)
- Even after this pass it can happen that you don't understand the paper
 - The subject is new to you
 - Unfamiliar with acronyms or terminology
 - The skipped details are way too relevant
 - The paper is written in a bad way...
- You can stop here eg. when you find this paper is interesting but outside your research scope

What if I don't get the paper?

- Well it depends on you...
- You can put it aside and hope you never need to look at it again
- Read it at another point in time
 - When you checked related work
 - Maybe the next morning with a fresh cup of coffee/tea
 - After talking to a colleague
 - Checking summaries
- Still continue with the 3rd pass

3rd pass

- Third Pass (> 1 hr.)
 - Fully understand the paper
 - Imagine you would like to re-implement the paper
 - This way you can identify hidden problems
 - Challenge every statement and assumption
 - Collect ideas for future work
 - For beginners it can take more than 4 hours
 - In the end you should be able to identify every problem
 - Even missing citations

Additional Tips

- Check online resources
 - Many people review popular papers
 - This can clear up questions
 - Check for YouTube Videos of the authors
 - Check if they have a poster for the paper
 - Search for YouTube videos of other people (eg. Yannic Kilcher, Two Minute Papers, AI coffee break)
 - Use AI tools to summarize the paper (and to ask questions about it)

Additional Tips

- Use paper libraries like Mendeley or Zotero to manage your paper collections
- Write down notes digital, eg. using Obsidian
- For literature review
 - Use your local library (USB)
 - Use Google Scholar
- Check if the paper is a pre-print or if it is peer-reviewed

